

High Performance Computing

Autumn, 2018

Lecture 1

Instructor

Prasun Ray

Teaching Fellow

Department of Mathematics

p.ray@imperial.ac.uk

Huxley 6M20

Office hours Tuesdays 5-6pm, MLC

Thursdays 4-5pm, MLC

(First office hour on Tuesday, 9/10)

Weekly schedule

- Lectures:
Tuesday, 2-3pm, Huxley 311
Thursday, 9-10am, Huxley 340
 - Labs:
Tuesday, 4-5pm, MLC (Huxley 414)
or
Wednesday, 10-11am, MLC
 - Only need to attend **one lab session**
- CDT students: alternate lab sessions have been arranged**

Syllabus

Lectures 1-2: Unix basics, version control with git/bitbucket

Lectures 3-6: Programming and scientific computing with Python

Lectures 7-10: Modular programming with Fortran, libraries, makefiles, coupling Fortran+Python

Syllabus

Lectures 11-14: Introduction to parallel computing and OpenMP

Lectures 15-16: Distributed memory computing with MPI, parallel libraries

Lectures 17-20: Basic computer architecture, cloud computing, cluster computing with Python and Spark

Assessment (tentative)

3 Programming assignments

HW1: Assigned 22/10, due 1/11 (**20%**)

HW2: Assigned 5/11, due 16/11 (**25%**)

HW3: Assigned 19/11, due 28/11 (**25%**)

1 Programming Project (**30%**)

Assigned 29/11, due 14/12

Submitting HW1 commits you to the course

CDT students: Will be contacted about assessment separately

Online material

- Main resource is course webpage:

<http://imperialm3c.bitbucket.io/>

- Slides will be available before every lecture

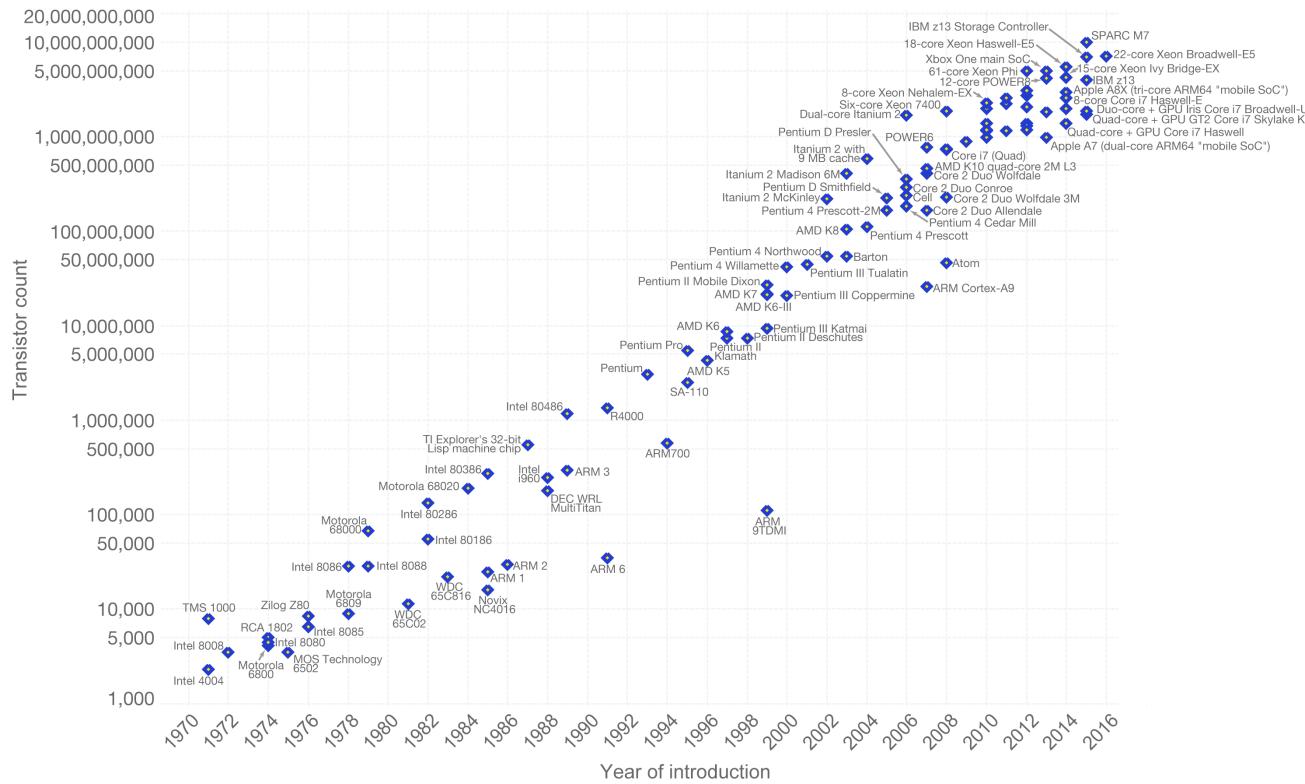
All course material will be available on course bitbucket page
(more on this later):

<https://bitbucket.org/imperialm3c/m3c2018>

Moore's law

Moore's Law – The number of transistors on integrated circuit chips (1971-2016) Our World in Data

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are strongly linked to Moore's law.



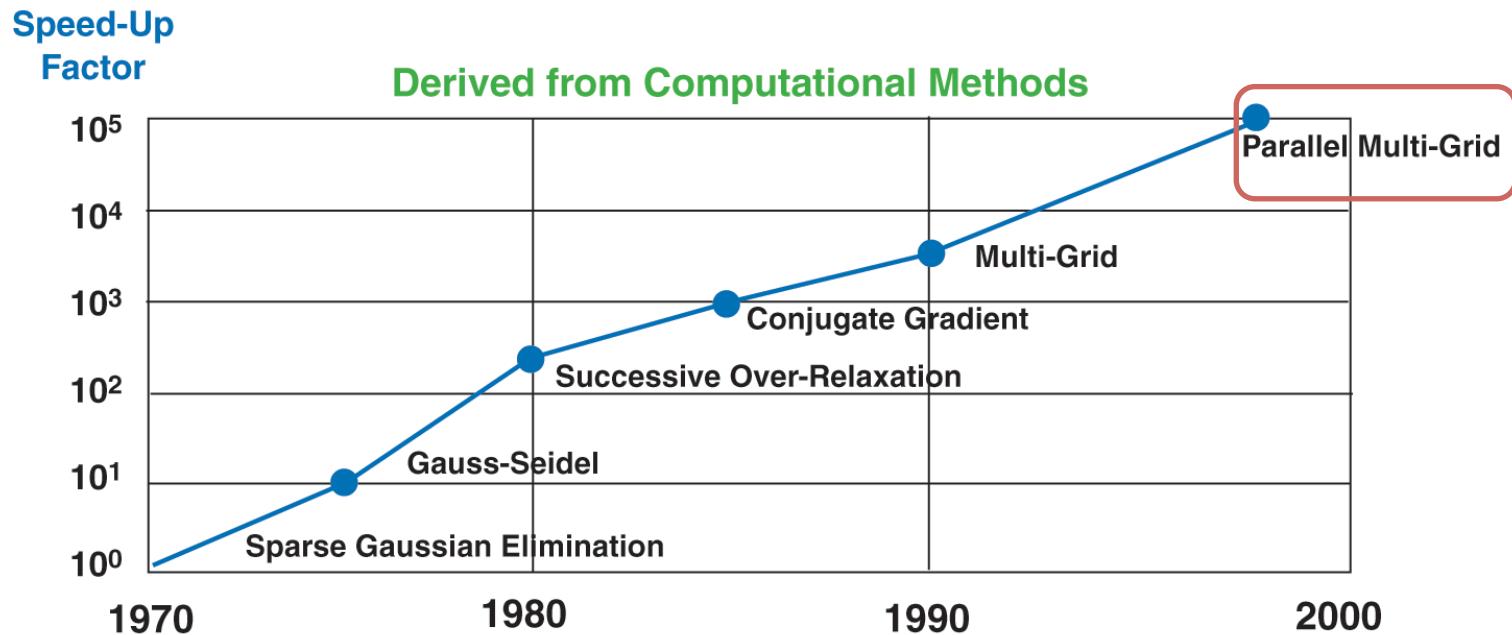
- Number of transistors on chip doubles every 2 years
 - From ~2005, multicore processors (why?)
 - 2 core, then 4, 6, 8, 16
 - Intel Xeon Phi: 60+ cores!

Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)

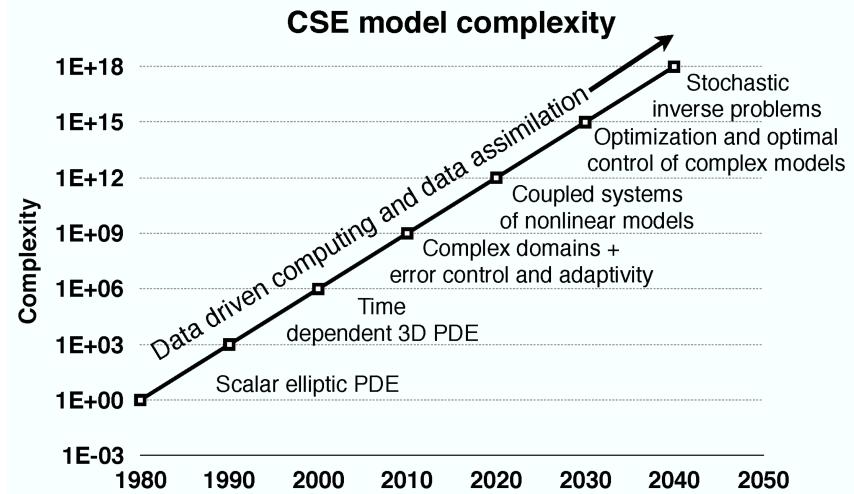
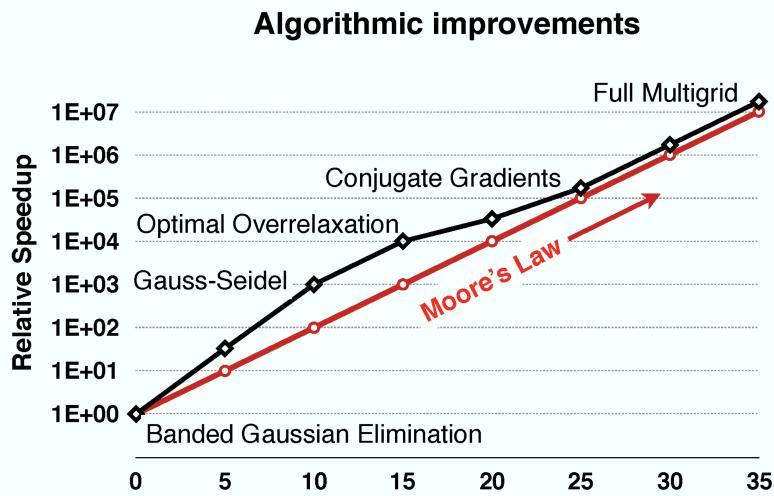
The data visualization is available at OurWorldInData.org. There you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser

Algorithms and hardware



Algorithms and hardware



High-end HPC

Rank	System	Cores	Rmax [TFlop/s]	Rpeak [TFlop/s]	Power [kW]
1	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM DOE/SC/Oak Ridge National Laboratory United States	2,282,544	122,300.0	187,659.3	8,806
2	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway , NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371
3	Sierra - IBM Power System S922LC, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM DOE/NNSA/LLNL United States	1,572,480	71,610.0	119,193.6	
4	Tianhe-2A - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, TH Express-2, Matrix-2000 , NUDT National Super Computer Center in Guangzhou China	4,981,760	61,444.5	100,678.7	18,482
5	AI Bridging Cloud Infrastructure [ABCi] - PRIMERGY CX2550 M4, Xeon Gold 6148 20C 2.4GHz, NVIDIA Tesla V100 SXM2, Infiniband EDR , Fujitsu National Institute of Advanced Industrial Science and Technology (AIST) Japan	391,680	19,880.0	32,576.6	1,649
6	Piz Daint - Cray XC50, Xeon E5-2690v3 12C 2.6GHz, Aries interconnect , NVIDIA Tesla P100 , Cray Inc. Swiss National Supercomputing Centre (CSCS) Switzerland	361,760	19,590.0	25,326.3	2,272
7	Titan - Cray XK7, Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x , Cray Inc. DOE/SC/Oak Ridge National Laboratory United States	560,640	17,590.0	27,112.5	8,209

Historically: cluster computing limited to national labs, research universities

But now...

Cluster computing is mainstream

Big data means big computers!



Cluster computing is mainstream



Google Cloud Platform



Course objective

- Cluster computing is not free!
- Important to:
 - choose right tools
 - use them effectively

This course provides foundation for “intelligent, informed” computing.

Software tools

Useful to classify tools as *scientific* or *general purpose*

Software tools

Useful to classify tools as *scientific* or *general purpose*

Examples:

Scientific	General purpose
Matlab	Python
Fortran	C++
R	Java

Software tools

Languages are *compiled* or *interpreted*

Software tools

Languages are *compiled* or *interpreted*

Compiled	Interpreted
Fortran	Python
C++	Matlab
Java	R

Software tools

This course:

Python: interpreted, general purpose

Fortran: compiled, scientific

Operating systems

Most HPC and scientific computing requires Unix (or Unix-like terminals)

Linux and Mac OS are built on Unix (and have terminal apps)

- Fairly straightforward to install course software

Windows:

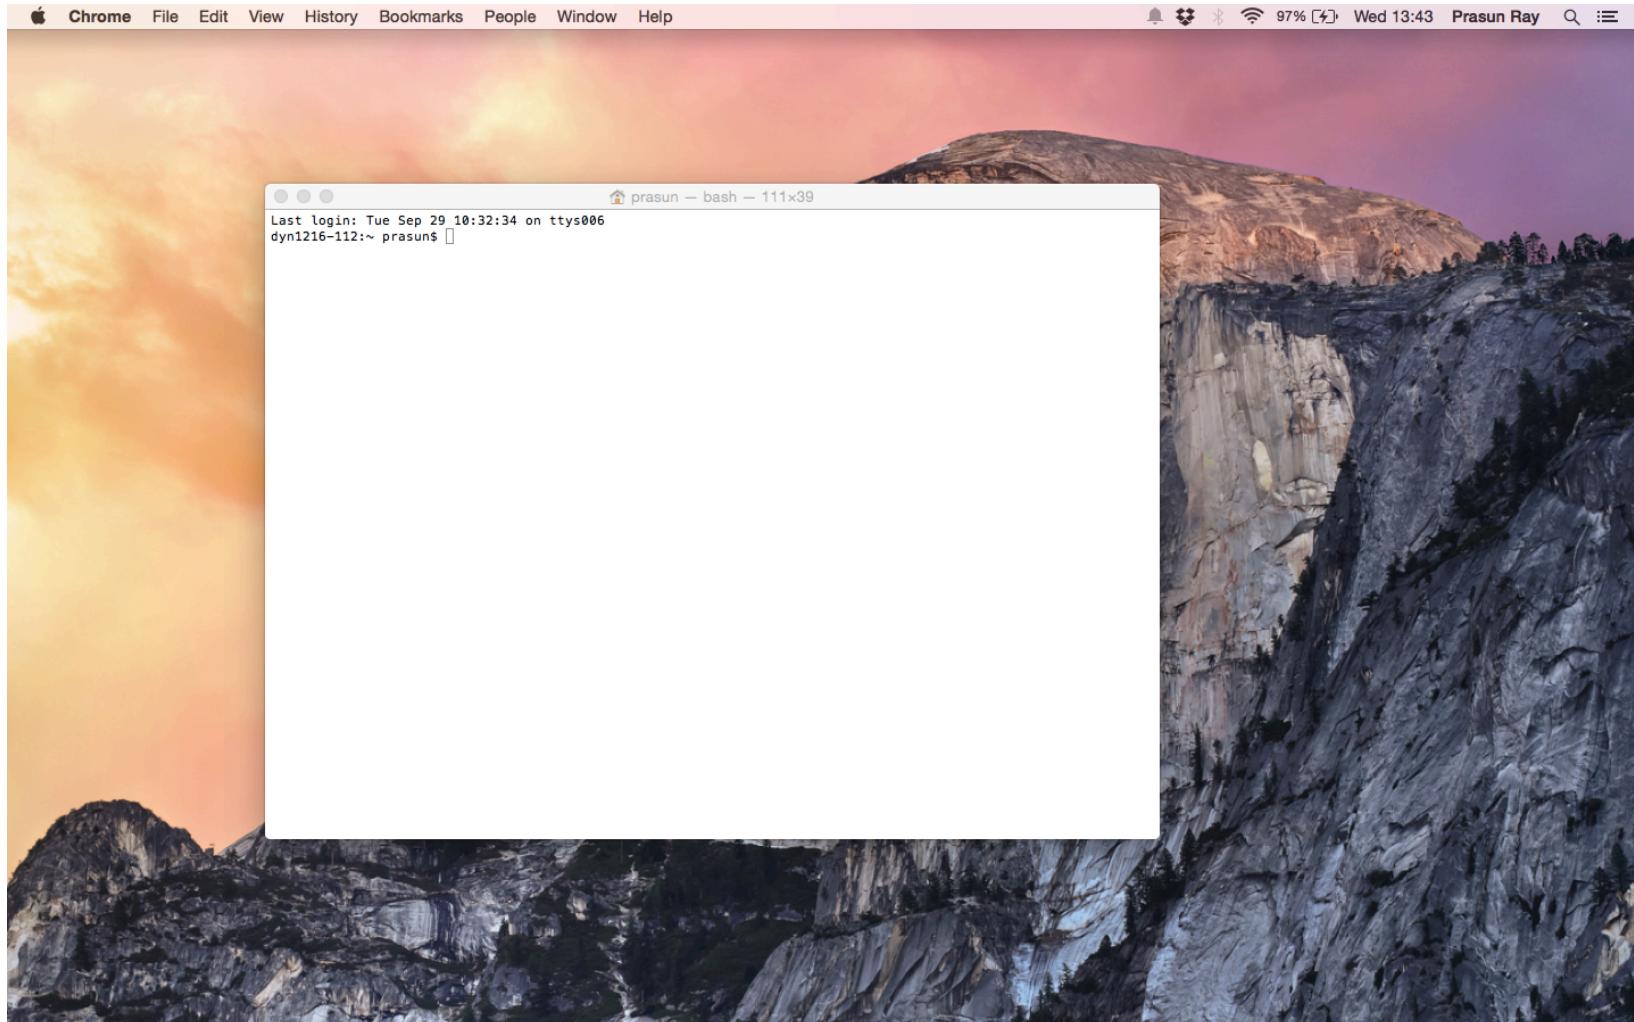
- Not well-suited for HPC
- Can get Unix terminal with cygwin
- For this course: Should install Linux virtual machine (VM) and install software within the VM
- MLC computers have Linux VMs installed (go try them out!)

Instructions for installing course software available online:

<http://imperialm3c.bitbucket.io/>

Unix terminal

Terminal on a mac:



12 Unix commands

Navigation:

pwd: print working directory (where am I?)

ls: list of directory contents (what is here?)

cd: change directory (let's go somewhere else)

```
$ pwd  
/Users/prasun/Documents/repos/m3c2017  
$  
$ ls  
Readme.md lectures  
$  
$ cd lectures  
$  
$ ls  
lecture1
```

12 Unix commands

Manipulate files and directories:

cp: Make copy of a file

mv: Move or rename a file

rm: Remove a file

rm -r: Remove directory and
all of its contents (dangerous!)

```
$ ls
Readme.md lectures
$ 
$ cp Readme.md Readme.md_copy
$ 
$ ls
Readme.md Readme.md_copy
lectures
$ 
$ mv Readme.md_copy
Readme.md_copy2
$ 
$ ls
Readme.md Readme.md_copy2
lectures
$ 
$ rm Readme.md_copy2
$ 
$ ls
Readme.md lectures
```

12 Unix commands

Info about contents of file:

cat: List contents of file

head -n: List first n lines

tail -n: list last n lines

grep: search within file for a string

```
$ cat example.txt
This is an example text file.
This is line 2.
This is line 3.
This is the last line.
$
$ head -1 example.txt
This is an example text file.
$
$ tail -2 example.txt
This is line 3.
This is the last line.
$
$ grep last example.txt
This is the last line.
```

12 Unix commands

Getting help:

man: manual page for a command

Try *man ls*. What does *ls -l* do? *ls -a*?

What if you don't know name of command?

https://en.wikipedia.org/wiki/List_of_Unix_commands

or google.

12 Unix commands

The 12 commands:

1. `pwd`
2. `ls`
3. `cd`
4. `cp`
5. `mv`
6. `rm`
7. `rm -r`
8. `cat`
9. `head -n`
10. `tail -n`
11. `grep`
12. `man`

This is “basic” Unix. Can do much more!

A little more Unix

Instead of outputting to screen, can output to file using “>”

```
$ ls  
example.txt  lecture1
```

```
$ grep last example.txt > output.txt
```

```
$ ls  
example.txt  lecture1  output.txt
```

```
$ cat output.txt  
This is the last line.
```

Lines in example.txt containing “last”
are written to output.txt

A little more Unix

Command can be executed sequentially (they can be “piped”) using “|”

```
$ head -2 example.txt | grep line > output.txt  
$  
$ cat output.txt  
This is line 2.
```



First two lines in example.txt are searched for the string “line” with results being written to output.txt

An example

You run optimization software that gives output that looks like:

```
INPUT:endgeom  
INPUT:azimuthal 9 0.1  
INPUT:polar 5  
INPUT:begin  
  
k-cactus is 1.402458
```

```
TIMING: Module: cpu      10.03 wall      10.04 Overall: cpu      29.00 wall      29.29  
=====
```

```
INPUT:EDIT 4
```

```
CALLING EDIT(INTERFACE_NO= 4)
```

```
INPUT:begin
```

```
*****  
INTERFACE 4 EIGENVALUE 1.402458 OVERALL MWd/t 0.0000E+00 BURNUP TIME 0.0000E+00  
DAYS  
*****
```

An example

We only care about the “k-cactus” values which appear several times.
How do we extract them?

```
INPUT:edgegeom  
INPUT:azimuthal 9 0.1  
INPUT:polar 5  
INPUT:begin
```

k-cactus is 1.402458

```
TIMING: Module: cpu      10.03 wall      10.04 Overall: cpu      29.00 wall      29.29  
=====  
=====  
INPUT:EDIT 4  
  
CALLING EDIT(INTERFACE_NO= 4)  
  
INPUT:begin  
*****  
INTERFACE 4 EIGENVALUE 1.402458 OVERALL MWd/t 0.0000E+00 BURNUP TIME 0.0000E+00  
DAYS  
=====
```

An example

Using grep:

```
$ grep cactus datafile.out
k-cactus is 1.402458
k-cactus is 1.386050
k-cactus is 1.377296
k-cactus is 1.352324
k-cactus is 1.328779
```

But what if we only want the numbers?

An example

Using grep:

```
$ grep cactus datafile.out
k-cactus is 1.402458
k-cactus is 1.386050
k-cactus is 1.377296
k-cactus is 1.352324
k-cactus is 1.328779
```

But what if we only want the numbers?

Use “cut”: \$ grep cactus datafile.out | cut -d s -f 3
1.402458
1.386050
1.377296
1.352324
1.328779

Questions: How do we store these numbers in a file? How do we find out what the flags after “cut” are doing?

What next?

- If you have your own laptop/desktop: start installing course software – see webpage for instructions
- If you don't, try out the virtual machines and these Unix commands on the MLC computers – go to software hub and search for “Linux – Maths” (and let me know if you have problems)
- Start working through introductory Python videos and exercises